

Bike-Share Case Study

This report provides the results as well as the step-by-step explanation of the data analysis performed for a bike-sharing case study. The data belongs to a bike-sharing company that has two kinds of users: annual members and casual riders. The goal of the case study was to identify how annual members and casual riders use the bikes differently in order to help the stake-holders decide whether or not to target converting casual riders into annual members in the next marketing campaign. The data used in this case study was from January-November 2023. Each month's data was stored in a csv file, and was downloaded from <https://divvy-tripdata.s3.amazonaws.com/index.html>, and Python was used to perform the analysis.

Dataset exploration & cleaning:

The code that was used to perform the data exploration can be found in the Jupyter Notebook [cleaning.ipynb](#). The main functions are explained below:

- *read_data*:

Here the .csv file for the bike rides of each month is read and stored into a dictionary called “data”. Each element in the dictionary has a key (the name of the month) and a value (the panada dataframe that holds the csv entries). This simplifies the access of the entries for each corresponding month, by using the month as the key (e.g. data[“May”] retrieves the dataframe that holds the entries from May). In Figure 1 we can see the first and last 5 entries of bike rides from May:

	ride_id	rideable_type	started_at	ended_at	start_station_name	start_station_id	end_station_name	end_station_id	start_lat	start_lng	end_lat	end_lng	member_casual
0	0D9FA920C3062031	electric_bike	2023-05-07 19:53:48	2023-05-07 19:58:32	Southport Ave & Belmont Ave	13229	NaN	NaN	41.939408	-87.663831	41.930000	-87.650000	member
1	92485E5FB5888ACD	electric_bike	2023-05-06 18:54:08	2023-05-06 19:03:35	Southport Ave & Belmont Ave	13229	NaN	NaN	41.939482	-87.663848	41.940000	-87.690000	member
2	FB144B3FC8300187	electric_bike	2023-05-21 00:40:21	2023-05-21 00:44:36	Halsted St & 21st St	13162	NaN	NaN	41.853793	-87.646719	41.860000	-87.650000	member
3	DDEB93BC2CE9AA77	classic_bike	2023-05-10 16:47:01	2023-05-10 16:59:52	Carpenter St & Huron St	13196	Damen Ave & Cortland St	13133	41.894556	-87.653449	41.915983	-87.677335	member
4	C07B70172FC92F59	classic_bike	2023-05-09 18:30:34	2023-05-09 18:39:28	Southport Ave & Clark St	TA1308000047	Southport Ave & Belmont Ave	13229	41.957081	-87.664199	41.939478	-87.663748	member
...
604822	48BDA26F34445546	electric_bike	2023-05-18 10:26:43	2023-05-18 10:48:00	Clark St & Elmdale Ave	KA1504000148	NaN	NaN	41.990876	-87.669721	42.000000	-87.660000	member
604823	573025E5EDE10DE1	electric_bike	2023-05-17 14:32:48	2023-05-17 14:45:37	State St & 33rd St	13216	NaN	NaN	41.834734	-87.625798	41.830000	-87.620000	member
604824	D88D48898C6FB63E	electric_bike	2023-05-17 07:59:29	2023-05-17 08:04:54	Columbus Dr & Randolph St	13263	NaN	NaN	41.884422	-87.619393	41.880000	-87.630000	member
604825	4692DCD2F87497F5	electric_bike	2023-05-18 08:34:48	2023-05-18 08:38:40	Public Rack - Karlov Ave & Lawrence Ave	1127.0	NaN	NaN	41.970000	-87.730000	41.970000	-87.740000	member
604826	6ACB7E383473D019	electric_bike	2023-05-29 21:16:58	2023-05-29 21:24:35	State St & 33rd St	13216	NaN	NaN	41.834715	-87.625764	41.840000	-87.650000	member

Figure 1: First and last 5 entries of bike rides from May (data[“May”])

From the entries we can see that the data for May consists of 13 columns: 1) ride id, 2) the type of bike, 3-4) date and time for the start and end of the ride, 5-12) the name, id, latitude and longitude of the start and end stations, and finally 13) whether the user was a casual rider or a member. In order to explore the dataset a bit further, the pandas method *nunique()* is used to return the number of unique values for each column in `data["May"]`. The output is shown in Figure 2:

Column Name:	NUnique
<code>ride_id</code>	604827
<code>rideable_type</code>	3
<code>started_at</code>	503683
<code>ended_at</code>	505259
<code>start_station_name</code>	1287
<code>start_station_id</code>	1250
<code>end_station_name</code>	1254
<code>end_station_id</code>	1210
<code>start_lat</code>	188591
<code>start_lng</code>	185410
<code>end_lat</code>	4759
<code>end_lng</code>	4762
<code>member_casual</code>	2

Figure 2: No. of unique values for every column in `data["May"]`

As is expected, the columns *rideable_type* and *member_casual* have a small number of unique values, whereas the rest of the columns do not. The pandas method *unique()* is then be used to find the unique values in these two columns:

- *rideable_type*: [electric_bike, classic_bike, docked_bike]
- *member_casual*: [member, casual]

- *count_entries*:

After the dataset is read into the dataframes, the method *count_entries* is used to collect further information about the dataset. It finds the number of entries per file as well as the number of columns. This is done in order to check that the data format is consistent across the different files. Next, the method calculates the total number of bike rides in the entire dataset. There is an option within the method to remove duplicates. Therefore, the method is first called with the remove duplicates option deactivated, in order to get a preliminary feel of the dataset. And then the method is called again with the remove duplicates option activated. The results are then written to output files which are shown in Figure 3.

Original_BikeRides			BikeRides_without_Duplicates		
Month	No Of Entries	No Of Cols	Month	No Of Entries	No Of Cols
January	190301	13	January	190301	13
February	190445	13	February	190445	13
March	258678	13	March	258678	13
April	426590	13	April	426590	13
May	604827	13	May	604827	13
June	719618	13	June	719618	13
July	767650	13	July	767650	13
August	771693	13	August	771693	13
September	666371	13	September	666371	13
October	537113	13	October	537113	13
November	362518	13	November	362518	13
Total:	5495804		Total:	5495804	
Average:	499618		Average:	499618	

Figure 3: No of entries and columns before and after removing duplicates

On the left is the result of running the method without removing duplicates, and on the right is the result after removing duplicates. We can see that all the files have the same number of columns, which is a good preliminary indicator of the consistency of data across the months. In total the dataset contains almost 5.5 Million entries, with an average of approximately 500,000 entries per month. From January to March the number of rides is relatively lower than the average, which is expected as these are cold months. This is confirmed by the peak in the number of rides during the Summer months June to August. The number of entries before and after removing duplicates is identical, therefore the original dataset did not have any duplicates.

- *check_NAN*:

Given that a brief look at the entries from May already showed a couple of NaN values, this method calculates the percentage of NaN values (Figure 1 *end_station_name* *end_station_id*). The number of null values for each column is calculated using the pandas function *isna()*. The percentage of null values for each column and month is shown in Figure 4. As we can see the columns *start_station_name*, *start_station_id*, *end_station_name*, *end_station_id* in every month have 13-17% null values. The columns *end_lat* and *end_long* have less than 1% null values.

- *clean_data*:

After exploring the dataset, we can see that the extractable information can be divided into information about the:

1. rider (casual/member)
2. bike (type of bike used)
3. ride (length, date, location)

NaN_Percentages													
Month	ride_id	rideable_type	started_at	ended_at	start_station_name	start_station_id	end_station_name	end_station_id	start_lat	start_lng	end_lat	end_lng	member_casual
January	0	0	0	0	14 %	14 %	14 %	14 %	0	0	< 1%	< 1%	0
February	0	0	0	0	13 %	13 %	14 %	14 %	0	0	< 1%	< 1%	0
March	0	0	0	0	13 %	13 %	14 %	14 %	0	0	< 1%	< 1%	0
April	0	0	0	0	14 %	14 %	16 %	16 %	0	0	< 1%	< 1%	0
May	0	0	0	0	14 %	14 %	15 %	15 %	0	0	< 1%	< 1%	0
June	0	0	0	0	16 %	16 %	17 %	17 %	0	0	< 1%	< 1%	0
July	0	0	0	0	16 %	16 %	16 %	16 %	0	0	< 1%	< 1%	0
August	0	0	0	0	15 %	15 %	16 %	16 %	0	0	< 1%	< 1%	0
September	0	0	0	0	15 %	15 %	16 %	16 %	0	0	< 1%	< 1%	0
October	0	0	0	0	15 %	15 %	16 %	16 %	0	0	< 1%	< 1%	0
November	0	0	0	0	15 %	15 %	15 %	15 %	0	0	< 1%	< 1%	0

Figure 4: Percentage of null values for each column across the various months

Since, the goal of the analysis is to find out how casual riders differ from members, it seems that the ride location is not of high relevance. It would have been important if for example the goal of the analysis was to find out whether more stations should be added and where to do so. However, in this case, and given that the locations seem to contain null values, it is safer to drop these columns. The column *ride_id* also does not provide any valuable information for the current analysis. Therefore, the relevant columns needed from this point onwards are the: *rideable_type*, *started_at*, *ended_at*, *member_casual*. So the first task performed by the method *clean_data* is to drop the columns that are no longer needed. Next, is data formatting. A quick check shows that the columns *started_at* and *ended_at* are stored as a strings of characters. Therefore, the second thing done by the method *clean_data* is to convert these columns into a datetime format and store the new dataset as *cleaned_data*. In order to ensure that the new dataset has the correct shape, the method *count_entries* is called once again and used to compare the cleaned dataset to the original one. The result is shown in Figure 5.

Original_BikeRides			BikeRides_Cleaned		
Month	No Of Entries	No Of Cols	Month	No Of Entries	No Of Cols
January	190301	13	January	190301	4
February	190445	13	February	190445	4
March	258678	13	March	258678	4
April	426590	13	April	426590	4
May	604827	13	May	604827	4
June	719618	13	June	719618	4
July	767650	13	July	767650	4
August	771693	13	August	771693	4
September	666371	13	September	666371	4
October	537113	13	October	537113	4
November	362518	13	November	362518	4
Total:	5495804		Total:	5495804	
Average:	499618		Average:	499618	

Figure 5: No of entries and columns before and after dropping the null values

As expected only the number of columns has changed (from 13 to 4), and the number of entries remains the same. As a final check, the method *check_NAN* is applied to the cleaned data, the result is as expected and is shown in Figure 6:

NaN_Percentages_Clean				
Month	rideable_type	started_at	ended_at	member_casual
January	0	0	0	0
February	0	0	0	0
March	0	0	0	0
April	0	0	0	0
May	0	0	0	0
June	0	0	0	0
July	0	0	0	0
August	0	0	0	0
September	0	0	0	0
October	0	0	0	0
November	0	0	0	0

Figure 6: Percentage of null values after cleaning

- *prepare_data*:

Now that the data is clean and in the correct format, we can extract the required information for analysis. First, the method *prepare_data* adds a new column *ride_length* which is the difference between the columns *ended_at* and *started_at*. Second,

Analysis: